

Recent Trends in Scientific Computing


John Stachurski

November 2023

Topics


- Traditional compiled languages
- Modern JIT compilers
- AI-driven scientific computing
- Where are we heading?
- Economic applications

Some trends:



Source: Stackoverflow Trends

Stack Overflow 2023 Developer Survey (50 languages)



— <https://survey.stackoverflow.co/2023/>

A review of some scientific computing environments

General purpose scientific computing environments:

1. Fortran & C / C++
2. MATLAB (\approx Python + NumPy)
3. Julia (\approx Python + Numba)
4. Python + Google JAX (\approx Python + PyTorch)

Fortran & C — static types and AOT compilers

Example. Suppose we want to compute the sequence

$$k_{t+1} = sk_t^\alpha + (1 - \delta)k_t$$

from some given k_0

Let's write a function in C that

1. implements the loop
2. returns the last k_t

```
#include <stdio.h>
```

```
#include <math.h>
```

```
int main() {  
    double k = 0.2;  
    double alpha = 0.4;  
    double s = 0.3;  
    double delta = 0.1;  
    int i;  
    int n = 1000;  
    for (i = 0; i < n; i++) {  
        k = s * pow(k, alpha) + (1 - delta) * k;  
    }  
    printf("k = %f\n", k);  
}
```

```
ϕ john on gz-precision .../treasury_2023 on β main
>> gcc solow.c -o out -lm
```

```
ϕ john on gz-precision .../treasury_2023 on β main
>> ./out
```

```
x = 6.240251
```

Pros

- fast

Cons

- time consuming to write
- lack of portability
- hard to debug
- hard to parallelize
- low interactivity

For comparison, the same operation in Python:

```
 $\alpha$  = 0.4  
s = 0.3  
 $\delta$  = 0.1  
n = 1_000  
k = 0.2  
  
for i in range(n-1):  
    k = s * k** $\alpha$  + (1 -  $\delta$ ) * k  
  
print(k)
```

Pros

- easy to write
- high portability
- easy to debug
- high interactivity

Cons

- slow

So how can we get


good execution speeds **and** high productivity / interactivity?

MATLAB

```
A = [2.0, -1.0  
      5.0, -0.5];
```

```
b = [0.5, 1.0]';
```

```
x = inv(A) * b
```



Python + NumPy


```
import numpy
```

```
A = ((2.0, -1.0),  
      (5.0, -0.5))
```

```
b = (0.5, 1.0)
```

```
A, b = np.array(A), np.array(b)
```

```
x = np.inv(A) @ b
```



Julia — rise of the JIT compilers

Can do MATLAB / NumPy style vectorized operations

$$A = \begin{bmatrix} 2.0 & -1.0 \\ 5.0 & -0.5 \end{bmatrix}$$
$$\mathbf{b} = [0.5 \quad 1.0]^T$$
$$x = \text{inv}(A) * b$$

But also has fast loops via an efficient JIT compiler

Example. Suppose, again, that we want to compute

$$k_{t+1} = sk_t^\alpha + (1 - \delta)k_t$$

from some given k_0

- Iterative, not easily vectorized

```
function solow(k0, α=0.4, δ=0.1, n=1_000)
    k = k0
    for i in 1:(n-1)
        k = s * k^α + (1 - δ) * k
    end
    return k
end

solow(0.2)
```

Julia accelerates `solow` at runtime via a JIT compiler

Python + Numba copy Julia

```
from numba import jit

@jit(nopython=True)
def solow(k0,  $\alpha=0.4$ ,  $\delta=0.1$ , n=1_000):
    k = k0
    for i in range(n-1):
        k = s * k** $\alpha$  + (1 -  $\delta$ ) * k
    return k

solow(0.2)
```

Runs at same speed as Julia / C / Fortran

AI-driven scientific computing

Key players

- TensorFlow, PyTorch
- Google JAX
- Mojo?

Examples.

- OpenAI uses PyTorch
- Google Bard uses Google JAX
- Apple Ajax uses Google JAX

Lightening introduction to deep learning

Supervised deep learning: find a good approximation to an unknown functional relationship

$$y = f(x)$$

- x is the input and y is the output

Examples.

- x = weather sensor data, y = max temp tomorrow
- x = income distribution, y = tax revenue next year
- x = unfinished sentence, y = next word

Training

Nonlinear regression: Take data set $(x_i, y_i)_{i=1}^n$ and solve

$$\min_{\theta} \ell(\theta) = \sum_{i=1}^n (y_i - \psi_{\theta}(x_i))^2 \quad \text{s.t.} \quad \theta \in \Theta$$


In the case of ANNs, we consider all ψ_{θ} having the form

$$\psi_{\theta} = \sigma \circ A_1 \circ \dots \circ \sigma \circ A_k$$


where

- $A_i x = W_i x + b_i$ is an affine map
- σ is a nonlinear “activation” function

Minimizing a smooth loss functions



Source: <https://danielkhv.com/>



Source: <https://losslandscape.com/gallery/>

Core elements

- automatic differentiation
- parallelization (CPUs / GPUs / TPUs)
- Compilers / JIT-compilers

```
import jax.numpy as jnp
from jax import grad, jit

def predict(params, x):
    for W, b in params:
        y = jnp.dot(W, x) + b
        x = jnp.tanh(y)
    return y

def loss(params, x, targets):
    preds = predict(params, x)
    return jnp.sum((preds - targets)**2)

grad_loss = jit(grad(loss))
```

Sample code

https://github.com/QuantEcon/treasury_2023